

```

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline

#from google.colab import drive
#drive.mount('/content/drive')

## Display all the columns of the dataframe
pd.set_option('display.max_columns',None)

df=pd.read_csv('/content/drive/MyDrive/PROJ-406-Capstone/telecom_customer_churn.csv')

```

```
df.head()
```

Customer ID	Unlimited Data	Contract	Paperless Billing	Payment Method	Monthly Charge	Total Charges	Total Refunds	Total Extra Data Charges	Total Long Distance Charges	Total Revenue	Customer Status	Churn Category	Churn Reason
No	Yes	One Year	Yes	Credit Card	65.6	593.30	0.00	0	381.51	974.81	Stayed	NaN	NaN
Yes	No	Month-to-Month	No	Credit Card	-4.0	542.40	38.33	10	96.21	610.28	Stayed	NaN	NaN
No	Yes	Month-to-Month	Yes	Bank Withdrawal	73.9	280.85	0.00	0	134.60	415.45	Churned	Competitor	Competitor had better devices
No	Yes	Month-to-Month	Yes	Bank Withdrawal	98.0	1237.85	0.00	0	361.66	1599.51	Churned	Dissatisfaction	Product dissatisfaction
No	Yes	Month-to-Month	Yes	Credit Card	83.9	267.40	0.00	0	22.14	289.54	Churned	Dissatisfaction	Network reliability

```
df.columns
```

```

Index(['Customer ID', 'Gender', 'Age', 'Married', 'Number of Dependents',
       'City', 'Zip Code', 'Latitude', 'Longitude', 'Number of Referrals',
       'Tenure in Months', 'Offer', 'Phone Service',
       'Avg Monthly Long Distance Charges', 'Multiple Lines',
       'Internet Service', 'Internet Type', 'Avg Monthly GB Download',
       'Online Security', 'Online Backup', 'Device Protection Plan',
       'Premium Tech Support', 'Streaming TV', 'Streaming Movies',
       'Streaming Music', 'Unlimited Data', 'Contract', 'Paperless Billing',
       'Payment Method', 'Monthly Charge', 'Total Charges', 'Total Refunds',
       'Total Extra Data Charges', 'Total Long Distance Charges',
       'Total Revenue', 'Customer Status', 'Churn Category', 'Churn Reason',
       'No. of Service'],
      dtype='object')

```

```
df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7043 entries, 0 to 7042
Data columns (total 39 columns):
 #   Column           Non-Null Count  Dtype  
 --- 
 0   Customer ID     7043 non-null    object  
 1   Gender          7043 non-null    object  
 2   Age             7043 non-null    int64  
 3   Married         7043 non-null    object  
 4   Number of Dependents  7043 non-null  int64  
 5   City            7043 non-null    object  
 6   Zip Code        7043 non-null    int64  
 7   Latitude        7043 non-null    float64 
 8   Longitude       7043 non-null    float64 
 9   Number of Referrals  7043 non-null  int64  
 10  Tenure in Months 7043 non-null    int64  
 11  Offer           3166 non-null    object  
 12  Phone Service   7043 non-null    object  
 13  Avg Monthly Long Distance Charges 6361 non-null  float64 

```

```

14  Multiple Lines          6361 non-null  object
15  Internet Service       7043 non-null  object
16  Internet Type          5517 non-null  object
17  Avg Monthly GB Download 5517 non-null  float64
18  Online Security         5517 non-null  object
19  Online Backup            5517 non-null  object
20  Device Protection Plan   5517 non-null  object
21  Premium Tech Support     5517 non-null  object
22  Streaming TV             5517 non-null  object
23  Streaming Movies          5517 non-null  object
24  Streaming Music           5517 non-null  object
25  Unlimited Data           5517 non-null  object
26  Contract                  7043 non-null  object
27  Paperless Billing         7043 non-null  object
28  Payment Method            7043 non-null  object
29  Monthly Charge            7043 non-null  float64
30  Total Charges              7043 non-null  float64
31  Total Refunds              7043 non-null  float64
32  Total Extra Data Charges    7043 non-null  int64
33  Total Long Distance Charges 7043 non-null  float64
34  Total Revenue              7043 non-null  float64
35  Customer Status            7043 non-null  object
36  Churn Category              1869 non-null  object
37  Churn Reason                1869 non-null  object
38  No. of Service              7043 non-null  int64
dtypes: float64(9), int64(7), object(23)
memory usage: 2.1+ MB

```

```
df.describe()
```

Magnitude	Number of Referrals	Tenure in Months	Avg Monthly Long Distance Charges		Avg Monthly GB Download	Monthly Charge	Total Charges	Total Refunds	Total Extra Data Charges	Total Long Distance Charges	Total Revenue
.000000	7043.000000	7043.000000	6361.000000	5517.000000	7043.000000	7043.000000	7043.000000	7043.000000	7043.000000	7043.000000	7043.000000
.756684	1.951867	32.386767	25.420517	26.189958	63.596131	2280.381264	1.962182	6.860713	749.099262	3034.379056	
.154425	3.001199	24.542061	14.200374	19.586585	31.204743	2266.220462	7.902614	25.104978	846.660055	2865.204542	
.301372	0.000000	1.000000	1.010000	2.000000	-10.000000	18.800000	0.000000	0.000000	0.000000	21.360000	
.788090	0.000000	9.000000	13.050000	13.000000	30.400000	400.150000	0.000000	0.000000	70.545000	605.610000	1
.595293	0.000000	29.000000	25.690000	21.000000	70.050000	1394.550000	0.000000	0.000000	401.440000	2108.640000	
.969795	3.000000	55.000000	37.680000	30.000000	89.750000	3786.600000	0.000000	0.000000	1191.100000	4801.145000	
.192901	11.000000	72.000000	49.990000	85.000000	118.750000	8684.800000	49.790000	150.000000	3564.720000	11979.340000	

Lets begin our data analysis by checkin ght edimension of our data frame.

```
df.shape
```

```
(7043, 39)
```

Now, lets check if we have any missing values in our data frame.

```
df.isnull()
```

	Customer ID	Gender	Age	Married	Number of Dependents	City	Zip Code	Latitude	Longitude	Number of Referrals	Tenure in Months	Offer	Phone Service	Avg Monthly Long Distance Charges	Multiple Lines
0	False	False	False	False	False	False	False	False	False	False	False	False	True	False	False
1	False	False	False	False	False	False	False	False	False	False	False	False	True	False	False
2	False	False	False	False	False	False	False	False	False	False	False	False	False	False	False
3	False	False	False	False	False	False	False	False	False	False	False	False	False	False	False
4	False	False	False	False	False	False	False	False	False	False	False	False	True	False	False
...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...
7038	False	False	False	False	False	False	False	False	False	False	False	False	False	False	False
7039	False	False	False	False	False	False	False	False	False	False	False	False	False	False	False
7040	False	False	False	False	False	False	False	False	False	False	False	False	False	False	False
7041	False	False	False	False	False	False	False	False	False	False	False	False	False	False	False
7042	False	False	False	False	False	False	False	False	False	False	False	False	True	False	True

7043 rows × 39 columns

df.isnull().sum()


```

Customer ID          0
Gender              0
Age                 0
Married             0
Number of Dependents 0
City                0
Zip Code            0
Latitude            0
Longitude           0
Number of Referrals 0
Tenure in Months    0
Offer               3877
Phone Service       0
Avg Monthly Long Distance Charges 682
Multiple Lines      682
Internet Service    0
Internet Type        1526
Avg Monthly GB Download 1526
Online Security     1526
Online Backup        1526
Device Protection Plan 1526
Premium Tech Support 1526
Streaming TV         1526
Streaming Movies     1526
Streaming Music      1526
Unlimited Data       1526
Contract             0
Paperless Billing    0
Payment Method       0
Monthly Charge       0
Total Charges        0
Total Refunds        0
Total Extra Data Charges 0
Total Long Distance Charges 0
Total Revenue        0
Customer Status      0
Churn Category       5174
Churn Reason          5174
No. of Service        0
dtype: int64

```

sns.heatmap(df.isnull(), yticklabels=False, cbar=False, cmap='viridis')

&lt;Axes: &gt;



```
# Lets check for the data types of all columns of our data frame
df.dtypes
```

Customer ID	object
Gender	object
Age	int64
Married	object
Number of Dependents	int64
City	object
Zip Code	int64
Latitude	float64
Longitude	float64
Number of Referrals	int64
Tenure in Months	int64
Offer	object
Phone Service	object
Avg Monthly Long Distance Charges	float64
Multiple Lines	object
Internet Service	object
Internet Type	object
Avg Monthly GB Download	float64
Online Security	object
Online Backup	object
Device Protection Plan	object
Premium Tech Support	object
Streaming TV	object
Streaming Movies	object
Streaming Music	object
Unlimited Data	object
Contract	object
Paperless Billing	object
Payment Method	object
Monthly Charge	float64
Total Charges	float64
Total Refunds	float64
Total Extra Data Charges	int64
Total Long Distance Charges	float64
Total Revenue	float64
Customer Status	object
Churn Category	object
Churn Reason	object
No. of Service	int64
dtype:	object

Lets find the current status of all the customer and easily visualize it using pie chart.

```

df_CurrentStatus = df['Customer Status'].value_counts().index
df_CurrentStatus

Index(['Stayed', 'Churned', 'Joined'], dtype='object', name='Customer Status')


df_CurrentStatus_values = df['Customer Status'].value_counts().values
df_CurrentStatus_values

array([4720, 1869, 454])

plt.pie(df_CurrentStatus_values[:3], labels = df_CurrentStatus[:3], autopct='%1.2f%')

([<matplotlib.patches.Wedge at 0x780144050190>,
 <matplotlib.patches.Wedge at 0x780144050070>,
 <matplotlib.patches.Wedge at 0x780144050d60>],
 [Text(-0.5604479811676993, 0.9465189170877941, 'Stayed'),
 Text(0.3586221713095951, -1.0398991000309556, 'Churned'),
 Text(1.0775211031566014, -0.22124256428676015, 'Joined')],
 [Text(-0.30569889881874507, 0.5162830456842513, '67.02%'),
 Text(0.1956120934415973, -0.5672176909259757, '26.54%'),
 Text(0.5877387835399643, -0.1206777623382328, '6.45%')])

```



Observation - It looks like majority of customers are staying with company but company is still loosing customer as 26.54% are churned and only 6.45% only are joining. In a way, company is effectively loosing its customer base by almost 20%.

```

df_ChurnCategory = df['Churn Category'].value_counts().index
df_ChurnCategory


Index(['Competitor', 'Dissatisfaction', 'Attitude', 'Price', 'Other'], dtype='object', name='Churn Category')

df_ChurnCategory_values = df['Churn Category'].value_counts().values
df_ChurnCategory_values

array([841, 321, 314, 211, 182])

plt.figure(figsize=(8, 6))
plt.bar(df_ChurnCategory, df_ChurnCategory_values, color='skyblue')
plt.xlabel('Churn Category')
plt.ylabel('Count')
plt.title('Churn Category Distribution')
plt.xticks(rotation=45) # Rotate x-axis labels for better readability
plt.tight_layout() # Adjust layout to prevent clipping of labels
plt.show()

```



```
colors = ['skyblue', 'salmon', 'lightgreen', 'orange']


# Plotting
plt.figure(figsize=(8, 6))
bars = plt.bar(df_ChurnCategory, df_ChurnCategory_values, color=colors)

# Add value labels to each bar
for bar, value in zip(bars, df_ChurnCategory_values):
    plt.text(bar.get_x() + bar.get_width() / 2, bar.get_height(), str(value),
             ha='center', va='bottom', color='black')

plt.xlabel('Churn Category')
plt.ylabel('Count')
plt.title('Churn Category Distribution')

# Customize x-axis labels rotation for better readability
plt.xticks(rotation=45)

plt.tight_layout() # Adjust layout to prevent clipping of labels
plt.show()
```



```
df.isna().sum()
```

Customer ID	0
Gender	0
Age	0
Married	0
Number of Dependents	0
City	0
Zip Code	0
Latitude	0
Longitude	0
Number of Referrals	0
Tenure in Months	0
Offer	3877
Phone Service	0
Avg Monthly Long Distance Charges	682
Multiple Lines	682
Internet Service	0
Internet Type	1526
Avg Monthly GB Download	1526
Online Security	1526
Online Backup	1526
Device Protection Plan	1526
Premium Tech Support	1526
Streaming TV	1526
Streaming Movies	1526
Streaming Music	1526
Unlimited Data	1526
Contract	0
Paperless Billing	0
Payment Method	0
Monthly Charge	0
Total Charges	0
Total Refunds	0
Total Extra Data Charges	0
Total Long Distance Charges	0
Total Revenue	0
Customer Status	0
Churn Category	5174
Churn Reason	5174
No. of Service	0

dtype: int64

Lets analyze how much our customers are paying per month and what is their tenure based on their status to have a better insights on their monthly expenses.

```
df.groupby('Customer Status')[['Monthly Charge', 'Tenure in Months']].agg(['min', 'max', 'mean'])
```

Customer Status	Monthly Charge			Tenure in Months			grid icon
	min	max	mean	min	max	mean	
Churned	-10.0	118.35	73.347592	1	72	17.979133	
Joined	-8.0	107.95	42.775991	1	3	1.720264	
Stayed	-10.0	118.75	61.737415	4	72	41.041525	

```
# I am creating one more dataframe to ensure that my original dataframe data is intact
dataset = df
```

dataset

	Customer ID	Gender	Age	Married	Number of Dependents	City	Zip Code	Latitude	Longitude	Number of Referrals	Tenure in Months	Offer	Phone Service	Avg Monthly Long Distance Charges	Mu:
0	0002-ORFBO	Female	37	Yes	0	Frazier Park	93225	34.827662	-118.999073	2	9	NaN	Yes	42.39	
1	0003-MKNFE	Male	46	No	0	Glendale	91206	34.162515	-118.203869	0	9	NaN	Yes	10.69	
2	0004-TLHLJ	Male	50	No	0	Costa Mesa	92627	33.645672	-117.922613	0	4	Offer E	Yes	33.65	
3	0011-IGKFF	Male	78	Yes	0	Martinez	94553	38.014457	-122.115432	1	13	Offer D	Yes	27.82	
4	0013-EXCHZ	Female	75	Yes	0	Camarillo	93010	34.227846	-119.079903	3	3	NaN	Yes	7.38	
...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	
7038	9987-LUTYD	Female	20	No	0	La Mesa	91941	32.759327	-116.997260	0	13	Offer D	Yes	46.68	
7039	9992-RRAMN	Male	40	Yes	0	Riverbank	95367	37.734971	-120.954271	1	22	Offer D	Yes	16.20	
7040	9992-UJOEL	Male	22	No	0	Elk	95432	39.108252	-123.645121	0	2	Offer E	Yes	18.62	
7041	9993-LHIEB	Male	21	Yes	0	Solana Beach	92075	33.001813	-117.263628	5	67	Offer A	Yes	2.12	
7042	9995-HOTOH	Male	36	Yes	0	Sierra City	96125	39.600599	-120.636358	1	63	NaN	No	NaN	

7043 rows × 39 columns

dataset.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7043 entries, 0 to 7042
Data columns (total 39 columns):
 #   Column          Non-Null Count  Dtype  
 --- 
 0   Customer ID    7043 non-null    object 
 1   Gender          7043 non-null    object 
 2   Age             7043 non-null    int64  
 3   Married         7043 non-null    object 
 4   Number of Dependents  7043 non-null    int64  
 5   City            7043 non-null    object 
 6   Zip Code        7043 non-null    int64  
 7   Latitude        7043 non-null    float64
 8   Longitude       7043 non-null    float64
 9   Number of Referrals  7043 non-null    int64  
 10  Tenure in Months 7043 non-null    int64  
 11  Offer           3166 non-null    object 
 12  Phone Service  7043 non-null    object 
```

```

13 Avg Monthly Long Distance Charges 6361 non-null float64
14 Multiple Lines 6361 non-null object
15 Internet Service 7043 non-null object
16 Internet Type 5517 non-null object
17 Avg Monthly GB Download 5517 non-null float64
18 Online Security 5517 non-null object
19 Online Backup 5517 non-null object
20 Device Protection Plan 5517 non-null object
21 Premium Tech Support 5517 non-null object
22 Streaming TV 5517 non-null object
23 Streaming Movies 5517 non-null object
24 Streaming Music 5517 non-null object
25 Unlimited Data 5517 non-null object
26 Contract 7043 non-null object
27 Paperless Billing 7043 non-null object
28 Payment Method 7043 non-null object
29 Monthly Charge 7043 non-null float64
30 Total Charges 7043 non-null float64
31 Total Refunds 7043 non-null float64
32 Total Extra Data Charges 7043 non-null int64
33 Total Long Distance Charges 7043 non-null float64
34 Total Revenue 7043 non-null float64
35 Customer Status 7043 non-null object
36 Churn Category 1869 non-null object
37 Churn Reason 1869 non-null object
38 No. of Service 7043 non-null int64
dtypes: float64(9), int64(7), object(23)
memory usage: 2.1+ MB

```

```

## Lets check the percentage of nan values present in each feature
## 1 -step make the list of features which has missing values
features_with_na=[features for features in dataset.columns if dataset[features].isnull().sum()>1]
## 2- step print the feature name and the percentage of missing values

```

```

for feature in features_with_na:
    print(feature, np.round(dataset[feature].isnull().mean(), 4), ' % missing values')

```

```

Offer 0.5505 % missing values
Avg Monthly Long Distance Charges 0.0968 % missing values
Multiple Lines 0.0968 % missing values
Internet Type 0.2167 % missing values
Avg Monthly GB Download 0.2167 % missing values
Online Security 0.2167 % missing values
Online Backup 0.2167 % missing values
Device Protection Plan 0.2167 % missing values
Premium Tech Support 0.2167 % missing values
Streaming TV 0.2167 % missing values
Streaming Movies 0.2167 % missing values
Streaming Music 0.2167 % missing values
Unlimited Data 0.2167 % missing values
Churn Category 0.7346 % missing values
Churn Reason 0.7346 % missing values

```

```

# Check for missing values in each column
missing_values = dataset.isnull().sum()

```

```

# Print the number of missing values in each column
print("Missing values in each column:")
for column, missing_count in missing_values.items():
    print(f"{column}: {missing_count}")

```

```

# Check for missing values in the entire dataset
total_missing = dataset.isnull().sum().sum()
print("\nTotal missing values in the dataset:", total_missing)

```

```

Missing values in each column:
Customer ID: 0
Gender: 0
Age: 0
Married: 0
Number of Dependents: 0
City: 0
Zip Code: 0
Latitude: 0
Longitude: 0
Number of Referrals: 0
Tenure in Months: 0
Offer: 3877
Phone Service: 0
Avg Monthly Long Distance Charges: 682
Multiple Lines: 682

```

```
Internet Service: 0
Internet Type: 1526
Avg Monthly GB Download: 1526
Online Security: 1526
Online Backup: 1526
Device Protection Plan: 1526
Premium Tech Support: 1526
Streaming TV: 1526
Streaming Movies: 1526
Streaming Music: 1526
Unlimited Data: 1526
Contract: 0
Paperless Billing: 0
Payment Method: 0
Monthly Charge: 0
Total Charges: 0
Total Refunds: 0
Total Extra Data Charges: 0
Total Long Distance Charges: 0
Total Revenue: 0
Customer Status: 0
Churn Category: 5174
Churn Reason: 5174
No. of Service: 0
```

Total missing values in the dataset: 30849

```
# Check for duplicate rows in the entire dataset
duplicate_rows = dataset[dataset.duplicated()]
print("Duplicate rows in the dataset:")
print(duplicate_rows)

Duplicate rows in the dataset:
Empty DataFrame
Columns: [Customer ID, Gender, Age, Married, Number of Dependents, City, Zip Code, Latitude, Longitude, Number of Referrals, Tenure in Months]
Index: []
```

```
# Import necessary libraries
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from sklearn.svm import SVC
from sklearn.metrics import accuracy_score, classification_report
from sklearn.model_selection import cross_val_score

# Define class variables
class_variable = "Customer Status"

dataset
```

	Customer ID	Gender	Age	Married	Number of Dependents	City	Zip Code	Latitude	Longitude	Number of Referrals	Tenure in Months	Offer	Phone Service	Avg Monthly Long Distance Charges	Mu:
0	0002-ORFBO	Female	37	Yes	0	Frazier Park	93225	34.827662	-118.999073	2	9	NaN	Yes	42.39	
1	0003-MKNFE	Male	46	No	0	Glendale	91206	34.162515	-118.203869	0	9	NaN	Yes	10.69	
2	0004-TLHLJ	Male	50	No	0	Costa Mesa	92627	33.645672	-117.922613	0	4	Offer E	Yes	33.65	
3	0011-IGKFF	Male	78	Yes	0	Martinez	94553	38.014457	-122.115432	1	13	Offer D	Yes	27.82	
4	0013-EXCHZ	Female	75	Yes	0	Camarillo	93010	34.227846	-119.079903	3	3	NaN	Yes	7.38	
...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...
7038	9987-LUTYD	Female	20	No	0	La Mesa	91941	32.759327	-116.997260	0	13	Offer D	Yes	46.68	
7039	9992-RRAMN	Male	40	Yes	0	Riverbank	95367	37.734971	-120.954271	1	22	Offer D	Yes	16.20	
7040	9992-UJOEL	Male	22	No	0	Elk	95432	39.108252	-123.645121	0	2	Offer E	Yes	18.62	
7041	9993-LHIEB	Male	21	Yes	0	Solana Beach	92075	33.001813	-117.263628	5	67	Offer A	Yes	2.12	
7042	9995-HOTOH	Male	36	Yes	0	Sierra City	96125	39.600599	-120.636358	1	63	NaN	No	NaN	

7043 rows × 39 columns

```
irrelevant_variables = ["Customer ID", "City", "Zip Code", "Latitude", "Longitude", "Phone Service", "Avg Monthly Long Distance Charges", "Multiple Internet Service", "Internet Type", "Avg Monthly GB Download", "Online Security", "Online Backup", "Device Protection Plan", "Premium Tech Support", "Streaming TV", "Streaming Movies", "Streaming Music", "Paperless Billing", "Payment Method"]
```

```
dataset.drop(irrelevant_variables, axis=1, inplace=True)
```

```
# Encode categorical variables
label_encoders = {}
for column in dataset.select_dtypes(include=["object"]).columns:
    label_encoders[column] = LabelEncoder()
    dataset[column] = label_encoders[column].fit_transform(dataset[column])
```

```
# Define features (X) and target variable (y)
X = dataset.drop(class_variable, axis=1)
y = dataset[class_variable]
```

```
# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

```
# Create and train classification models
models = {
    "Logistic Regression": LogisticRegression(),
    "Random Forest": RandomForestClassifier(),
    "Support Vector Machine": SVC()
}
```

```
for name, model in models.items():
    model.fit(X_train, y_train)
```

```
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_logistic.py:458: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. OF ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (`max_iter`) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

`n_iter_i = _check_optimize_result()`

```
# Evaluate model
y_pred = model.predict(X_test)
accuracy = accuracy_score(y_test, y_pred)
cv_score = cross_val_score(model, X, y, cv=5).mean()
print(f"Model: {name}")
print(f"Accuracy: {accuracy}")
print(f"Cross-Validation Score: {cv_score}")
print(f"Classification Report:\n{classification_report(y_test, y_pred)}")
print("-----")
```

Model: Support Vector Machine  
Accuracy: 0.7111426543647977  
Cross-Validation Score: 0.7238395420672301

Classification Report:

	precision	recall	f1-score	support
0	0.45	0.39	0.41	373
1	0.00	0.00	0.00	97
2	0.79	0.91	0.85	939
accuracy			0.71	1409
macro avg	0.41	0.43	0.42	1409
weighted avg	0.64	0.71	0.67	1409

-----

/usr/local/lib/python3.10/dist-packages/sklearn/metrics/\_classification.py:1344: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0. This will change in version 0.20.  
\_warn\_prf(average, modifier, msg\_start, len(result))  
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/\_classification.py:1344: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0. This will change in version 0.20.  
\_warn\_prf(average, modifier, msg\_start, len(result))  
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/\_classification.py:1344: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0. This will change in version 0.20.  
\_warn\_prf(average, modifier, msg\_start, len(result))

```
# Increase max_iter for Logistic Regression
models = {
    "Logistic Regression": LogisticRegression(max_iter=1000), # Adjust max_iter as needed
    "Random Forest": RandomForestClassifier(),
    "Support Vector Machine": SVC()
}

for name, model in models.items():
    model.fit(X_train, y_train)

# Evaluate model
y_pred = model.predict(X_test)
accuracy = accuracy_score(y_test, y_pred)
cv_score = cross_val_score(model, X, y, cv=5).mean()
print(f"Model: {name}")
print(f"Accuracy: {accuracy}")
print(f"Cross-Validation Score: {cv_score}")
print(f"Classification Report:\n{classification_report(y_test, y_pred)}")
print("-----")
```

```

Model: Random Forest
Accuracy: 1.0
Cross-Validation Score: 1.0
Classification Report:
precision    recall   f1-score   support
0            1.00     1.00     1.00      373
1            1.00     1.00     1.00       97
2            1.00     1.00     1.00     939

accuracy          1.00      1.00      1.00     1409
macro avg       1.00     1.00     1.00     1409
weighted avg    1.00     1.00     1.00     1409
-----
```

```

Model: Support Vector Machine
Accuracy: 0.7111426543647977
Cross-Validation Score: 0.7238395420672301
Classification Report:
precision    recall   f1-score   support
0            0.45     0.39     0.41      373
1            0.00     0.00     0.00       97
2            0.79     0.91     0.85     939

accuracy          0.71      0.71      0.71     1409
macro avg       0.41     0.43     0.42     1409
weighted avg    0.64     0.71     0.67     1409
-----
```

```

/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1344: UndefinedMetricWarning: Precision and F-score are ill-defined and hence are ignored in label 2
    _warn_prf(average, modifier, msg_start, len(result))
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1344: UndefinedMetricWarning: Precision and F-score are ill-defined and hence are ignored in label 2
    _warn_prf(average, modifier, msg_start, len(result))
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1344: UndefinedMetricWarning: Precision and F-score are ill-defined and hence are ignored in label 2
    warn(nrf(average, modifier, msg_start, len(result)))
-----
```

```

from sklearn.model_selection import GridSearchCV

# Define SVM with hyperparameter grid for tuning
svm_params = {
    'C': [0.1, 1, 10],
    'kernel': ['linear', 'rbf', 'poly']
}
svm_model = SVC()
svm_grid_search = GridSearchCV(svm_model, svm_params, cv=5)
svm_grid_search.fit(X_train, y_train)

# Best parameters for SVM
print("Best parameters for SVM:", svm_grid_search.best_params_)

# Evaluate best SVM model
best_svm_model = svm_grid_search.best_estimator_
best_svm_model.fit(X_train, y_train)
y_pred_svm = best_svm_model.predict(X_test)
accuracy_svm = accuracy_score(y_test, y_pred_svm)
cv_score_svm = cross_val_score(best_svm_model, X, y, cv=5).mean()

print("Model: Support Vector Machine (Tuned)")
print(f"Accuracy: {accuracy_svm}")
print(f"Cross-Validation Score: {cv_score_svm}")
print(f"Classification Report:\n{classification_report(y_test, y_pred_svm)}")
print("-----")
```

```

Best parameters for SVM: {'C': 10, 'kernel': 'linear'}
Model: Support Vector Machine (Tuned)
Accuracy: 1.0
Cross-Validation Score: 0.9998579545454545
Classification Report:
precision    recall   f1-score   support
0            1.00     1.00     1.00      373
1            1.00     1.00     1.00       97
2            1.00     1.00     1.00     939

accuracy          1.00      1.00      1.00     1409
macro avg       1.00     1.00     1.00     1409
weighted avg    1.00     1.00     1.00     1409
-----
```

```

import numpy as np

# Define metrics
metrics = ['Precision', 'Recall', 'F1-score']
classes = ['Churned', 'Stayed']

# Define values for Logistic Regression
lr_values = [[0.85, 0.91, 0.88], [1.0, 0.94, 0.97]]

# Define values for Random Forest
rf_values = [[1.0, 1.0, 1.0], [1.0, 1.0, 1.0]]

# Define values for Support Vector Machine (Tuned)
svm_values = [[1.0, 1.0, 1.0], [1.0, 1.0, 1.0]]

# Plot metrics for each class
for i, metric in enumerate(metrics):
    plt.figure(figsize=(10, 5))

    # Combine values for each class
    lr_class_values = np.array(lr_values)[:, i]
    rf_class_values = np.array(rf_values)[:, i]
    svm_class_values = np.array(svm_values)[:, i]

    # Plot bar chart
    x = np.arange(len(classes))
    width = 0.2
    plt.bar(x - width, lr_class_values, width, label='Logistic Regression', color='b')
    plt.bar(x, rf_class_values, width, label='Random Forest', color='g')
    plt.bar(x + width, svm_class_values, width, label='SVM (Tuned)', color='r')

    plt.xlabel('Class')
    plt.ylabel(metric)
    plt.title(f'{metric} for Churned and Stayed Customers')
    plt.xticks(x, classes)
    plt.legend()
    plt.show()

```

